

Solution:

(a)

Denote $MWVC(v, j)$ = minimum weight of a vertex cover in the subtree rooted at v using exactly j vertices, with the convention that if no such cover exists then $MWVC(v, j) = +\infty$.

Base Case: For a leaf v ,

$$MWVC(v, 0) = 0, \quad MWVC(v, 1) = w(v), \quad MWVC(v, j) = +\infty \quad \text{for } j > 1.$$

Recursive Step: For an internal vertex v with children u_1, u_2, \dots, u_m ,

$$MWVC(v, j) = \min \left\{ \begin{array}{l} \min_{j_1 + \dots + j_m = j-1} \left[w(v) + \sum_{i=1}^m MWVC(u_i, j_i) \right], \\ \min_{\substack{j_1 + \dots + j_m = j \\ j_i \geq 1 \forall i}} \left[\sum_{i=1}^m MWVC(u_i, j_i) \right] \end{array} \right\}.$$

The first term corresponds to the case where v is included in the cover, and the second term corresponds to the case where v is not included (forcing every child u_i to be included).

Final Answer: Since we want a vertex cover of size *at most* k , the answer for the tree (with root r) is

$$\min_{0 \leq j \leq k} MWVC(r, j).$$

Time Analysis

For each vertex v , we compute $MWVC(v, j)$ for $j = 0, \dots, k$. By merging the DP tables of the children in $O(k^2)$ time per merge, if v has at most d children, the processing time per vertex is $O(d \cdot k^2)$. Thus, over all n vertices, the total time is

$$O(n \cdot d \cdot k^2).$$

If d is a constant, this simplifies to $O(n \cdot k^2)$.

- (b) (i) Let $T = (V, E)$ be a tree with an arbitrary root r . We define two functions for every vertex v :

$f(v)$ = the number of vertex covers of the subtree rooted at v that include v ,
 $g(v)$ = the number of vertex covers of the subtree rooted at v that exclude v .

Thus, the total number of vertex covers for the subtree rooted at v is:

$$C(v) = f(v) + g(v).$$

Base Case: If v is a leaf (i.e. v has no children), then the vertex cover on v can either include v or not, giving:

$$f(v) = 1, \quad g(v) = 1, \quad \text{so } C(v) = 2.$$

Recursive Step: For an internal vertex v with children u_1, u_2, \dots, u_m , we have two cases:

- **Case 1: v is included in the cover.** Since v covers all edges (v, u_i) , each child u_i can be either in or out of the cover. Hence, the number of vertex covers in this case is:

$$f(v) = \prod_{i=1}^m (f(u_i) + g(u_i)).$$

- **Case 2: v is not included in the cover.** If v is excluded, then to cover the edge (v, u_i) each child u_i *must* be included in the cover. Thus:

$$g(v) = \prod_{i=1}^m f(u_i).$$

Final Answer: The total number of distinct vertex covers for the entire tree T is given by:

$$C(r) = f(r) + g(r).$$

Time Analysis

Each vertex $v \in V$ is processed exactly once in a postorder traversal of the tree. For each vertex, we perform a constant amount of work (multiplying the results from its children). Hence, the overall time complexity of the algorithm is $O(|V|)$.

- (ii) We consider a path P_n on n nodes. For a single node (i.e. $n = 1$) the vertex cover can be either the empty set (which is valid since there are no edges) or the singleton set containing the node. Thus, the base case is:

$$DP[1] = 2.$$

For convenience, we also define a base case for the empty path (no nodes) as:

$$DP[0] = 1.$$

Now, for $n \geq 2$, consider the rightmost node (say node n):

- **Case 1:** Node n is *included* in the vertex cover. In this case, the remaining $n - 1$ nodes can be covered in any valid way. This contributes $DP[n - 1]$ covers.
- **Case 2:** Node n is *not included*. Then to cover the edge $(n - 1, n)$, node $n - 1$ *must* be included. In this case, the first $n - 2$ nodes can be covered in any valid way. This contributes $DP[n - 2]$ covers.

Thus, the recurrence is:

$$DP[n] = DP[n - 1] + DP[n - 2], \quad \text{for } n \geq 2.$$

This recurrence is exactly that of the Fibonacci sequence. To see the correspondence, define a Fibonacci sequence with the following initial conditions:

$$F[0] = 1, \quad F[1] = 1, \quad \text{and } F[n] = F[n-1] + F[n-2] \text{ for } n \geq 2.$$

Then, our recurrence satisfies:

$$DP[0] = F[1] = 1, \quad DP[1] = 2 = F[2],$$

and by induction one can show that

$$DP[n] = F[n+1].$$

Fitting in a 64-bit Word: The Fibonacci numbers grow exponentially; indeed, using Binet's formula we have

$$F[n] \approx \frac{\varphi^n}{\sqrt{5}},$$

where $\varphi \approx 1.618$. For $n = 501$, we have

$$DP[500] = F[501] \approx \frac{\varphi^{501}}{\sqrt{5}}.$$

Taking logarithms (base 10):

$$\log_{10} F[501] \approx 501 \cdot \log_{10} \varphi - \frac{1}{2} \log_{10} 5 \approx 501 \cdot 0.208987 - 0.349 \approx 104.44.$$

Thus, $F[501]$ has roughly 105 decimal digits. A 64-bit integer can hold values up to roughly $2^{63} \approx 9.22 \times 10^{18}$, which is only about 19 decimal digits. Therefore, the answer for a path with $n = 500$ does *not* fit in a 64-bit integer.

- (iii) **Implementation:** Use multi-precision arithmetic by representing large integers as arrays of 64-bit words. Each addition is performed word-by-word with carry propagation. For example:

- Use an array (or vector) where each entry stores a 64-bit word representing a “digit” in a fixed base (2^{64}).
- Implement addition on these arrays by performing word-level additions with carry propagation.

Running Time: There are n steps, and each addition takes $O(n)$ machine-word operations (since numbers have $O(n)$ bits), yielding an overall running time of $O(n^2)$.

■

Solution:

(a) The algorithm consists of two main steps:

- **DFA Construction:** We construct a deterministic finite automaton (DFA) $D = (Q, \Sigma, \delta, q_0, F)$ that accepts exactly those strings which contain the substring “374”. The alphabet Σ is $\{0, 1, \dots, 9\}$. The states of D represent the progress in matching the pattern “374”. For example, let

$$Q = \{q_0, q_1, q_2, q_3\},$$

where:

- q_0 is the start state (no characters matched),
- q_1 indicates that 3 has been matched,
- q_2 indicates that 37 has been matched,
- q_3 is the accepting state meaning that 374 has been matched.

The transition function δ is defined appropriately to update the match progress based on the next digit.

- **Product Graph and DFS:** We now construct the product graph $G' = (V', E')$ where:

$$V' = V \times Q.$$

There is an edge from (u, q) to (v, q') in E' if and only if $(u, v) \in E$ in G and $q' = \delta(q, \ell(v))$. We then perform a DFS (or BFS) starting from the initial state (s, q_0) in the product graph and check if we can reach a vertex of the form (t, q_3) (i.e., t paired with the accepting state q_3).

Algorithm 1 Find an s - t walk with substring 374

```

1: Input: Graph  $G = (V, E)$ , labeling function  $\ell : V \rightarrow \{0, 1, \dots, 9\}$ , vertices  $s, t$ 
2: Output: true if there exists an  $s$ - $t$  walk  $W$  in  $G$  with  $\sigma(W)$  containing 374, else false
3: Construct DFA  $D = (Q, \Sigma, \delta, q_0, F)$  where  $Q = \{q_0, q_1, q_2, q_3\}$  and  $F = \{q_3\}$ .
4: Initialize stack  $S \leftarrow \{(s, q_0)\}$ 
5: Initialize visited set  $Visited \leftarrow \{(s, q_0)\}$ 
6: while  $S$  is not empty do
7:    $(u, q) \leftarrow \text{pop from } S$ 
8:   if  $u = t$  and  $q = q_3$  then
9:     return true
10:  end if
11:  for each edge  $(u, v) \in E$  do
12:     $q' \leftarrow \delta(q, \ell(v))$ 
13:    if  $(v, q') \notin Visited$  then
14:      Add  $(v, q')$  to  $Visited$ 
15:      Push  $(v, q')$  onto  $S$ 
16:    end if
17:  end for
18: end while
19: return false

```

Time Analysis

- The DFA D has a constant number of states (4 in this case) and can be constructed in constant time.
- The product graph G' has $O(|V|)$ vertices (specifically, $|V| \times 4$) and each edge $(u, v) \in E$ creates at most 4 transitions in G' . Thus, the total number of transitions is $O(|E|)$.
- The DFS on the product graph runs in $O(|V'| + |E'|)$ time, which is $O(|V| + |E|)$.

Thus, the overall algorithm runs in $O(|V| + |E|)$ time.

- (b) We utilize the same DFA and product graph G' from (a) if there exists an s - t walk W in G with $\sigma(W)$ containing 374.

The algorithm proceeds in the following steps:

- **Reverse Reachability Search:** We wish to know for which vertices v there exists a walk from (v, q_0) to (t, q_3) in G' . Instead of searching from every possible v , we reverse the product graph and perform a DFS (or BFS) starting from the target state (t, q_3) . Let R be the set of all product states that can reach (t, q_3) in the reverse graph.
- **Output the Answer:** The final answer is the set of all vertices $v \in V$ for which (v, q_0) is in R .

Algorithm 2 Output set of vertices with a v - t walk containing substring 374

```

1: Input: Graph  $G = (V, E)$ , labeling function  $\ell : V \rightarrow \{0, 1, \dots, 9\}$ , vertex  $t$ 
2: Output: Set  $S \subseteq V$  of all vertices  $v$  such that there exists a  $v$ - $t$  walk  $W_v$  with  $\sigma(W_v)$ 
   containing 374

3: /* Step 1: Construct DFA for substring 374 */
4: Define  $Q = \{q_0, q_1, q_2, q_3\}$  with start state  $q_0$  and accepting state  $F = \{q_3\}$ 
5: Define transition function  $\delta : Q \times \{0, 1, \dots, 9\} \rightarrow Q$  (details omitted for brevity)

6: /* Step 2: Build the product graph implicitly */
7: Let  $V' = \{(v, q) : v \in V, q \in Q\}$ 
8: /* We do not explicitly construct  $E'$ ; instead we use  $\delta$  on the fly */

9: /* Step 3: Reverse reachability from  $(t, q_3)$  */
10: Initialize stack  $S \leftarrow \{(t, q_3)\}$ 
11: Initialize visited set  $Visited \leftarrow \{(t, q_3)\}$ 

12: while  $S$  is not empty do
13:   Pop  $(u, q)$  from  $S$ 
14:   for each edge  $(v, u) \in E$  do                                 $\triangleright$  Reverse edge:  $u$  is reached from  $v$ 
15:     Compute  $q'$  such that  $\delta(q', \ell(u)) = q$    $\triangleright$  (Determine predecessor DFA state  $q'$  given
       that reading  $\ell(u)$  leads to  $q$ .)
16:     if  $(v, q') \notin Visited$  then
17:       Add  $(v, q')$  to  $Visited$ 
18:       Push  $(v, q')$  onto  $S$ 
19:     end if
20:   end for
21: end while

22: /* Step 4: Collect all vertices  $v$  with  $(v, q_0)$  in  $Visited$  */
23:  $Result \leftarrow \{v \in V : (v, q_0) \in Visited\}$ 
24: return  $Result$ 

```

Time Analysis

- **DFA Construction:** The DFA has a constant number of states (4) and is constructed in $O(1)$ time.
 - **Product Graph:** The product graph G' has $|V| \times 4 = O(|V|)$ vertices. Each edge $(u, v) \in E$ yields a transition in the product graph. Thus, there are $O(|E|)$ transitions.
 - **Reverse DFS:** The reverse reachability search traverses at most all product vertices and edges, in time $O(|V'| + |E'|) = O(|V| + |E|)$.
 - **Overall Complexity:** The algorithm runs in $O(|V| + |E|)$ time.
- (c) We construct a new (configuration) graph whose vertices are triples (u, v, w) with u, v, w being distinct vertices of G . There is an edge from (a, b, c) to (a', b', c') if and only if

$$a' \in N(a), \quad b' \in N(b), \quad c' \in N(c),$$

and a', b', c' are all distinct. Here, $N(a)$ denotes the *neighborhood* of the vertex a in G , i.e., the set of all vertices that are adjacent to a in G .

We then require that the tokens only trigger the printouts at the correct times by dividing the motion into three phases:

- **Phase 1:** From (s_1, s_2, s_3) to $(1, 7, 3)$ while *avoiding* $(3, 7, 4)$.
- **Phase 2:** From $(1, 7, 3)$ to $(3, 7, 4)$ while *avoiding* any visit to $(1, 7, 3)$ or $(3, 7, 4)$ except at the endpoints.
- **Phase 3:** From $(3, 7, 4)$ to (t_1, t_2, t_3) while avoiding any visit to $(1, 7, 3)$ or $(3, 7, 4)$.

A path with these properties will produce the printed string 173 upon entering phase 1's target state, then 374 upon finishing phase 2, and no further printed substrings during phase 3.

Algorithm 3 Check for a Valid Sequence Producing 173374

```

1: procedure VALIDPATH( $G, s = (s_1, s_2, s_3), t = (t_1, t_2, t_3)$ )
2:   Let  $S = \{(u, v, w) \mid u, v, w \text{ distinct vertices of } G\}$ 
3:   Construct configuration graph  $H = (S, E)$  where

      Edge  $((a, b, c), (a', b', c')) \in E \iff a' \in N(a), b' \in N(b), c' \in N(c)$  and  $a', b', c'$  are distinct.

4:   Define special states:
            $p = (1, 7, 3)$  and  $q = (3, 7, 4)$ .

5:   // Phase 1: Find a path in  $H$  from  $s$  to  $p$  avoiding  $q$ .
6:   if BFS( $H, s, p$ , forbidden= $\{q\}$ ) returns false then
7:     return false
8:   end if
9:   // Phase 2: Find a path in  $H$  from  $p$  to  $q$  with no intermediate occurrence of  $p$  or  $q$ .
10:  if BFS( $H, p, q$ , forbidden= $\{p, q\} \setminus \{p, q\}$ ) returns false then
11:    return false
12:  end if
13:  // Phase 3: Find a path in  $H$  from  $q$  to  $t$  avoiding  $\{p, q\}$ .
14:  if BFS( $H, q, t$ , forbidden= $\{p, q\}$ ) returns false then
15:    return false
16:  end if
17:  return true
18: end procedure

```

Here, BFS is a breadth-first search procedure modified so that it does not traverse any vertex in the given `forbidden` set (except when the start or target is itself in that set, as allowed in the endpoints for phase 2).

Let n be the number of vertices in G . In the configuration graph H :

- The number of vertices is $V = O(n^3)$ (since each state is a triple of distinct vertices).
- Each vertex has at most $O(\Delta^3)$ neighbors, where Δ is the maximum degree of G . Hence, the number of edges is

$$E = O(n^3 \cdot \Delta^3).$$

Since a BFS runs in $O(V + E)$ time, one BFS takes

$$O(n^3 + n^3 \Delta^3) = O(n^3 \Delta^3).$$

As we perform a constant number (three) of BFS operations, the overall time complexity of the algorithm is

$$O(n^3 \Delta^3).$$

If Δ is bounded by a constant, this further simplifies to $O(n^3)$.

■